

# 强化学习及其应用

Reinforcement Learning and Its Applications

## 第三章 值估计

Value Evaluation

授课人：周晓飞

zhouxiaofei@iie.ac.cn

2018-6-26

## 第三章 值估计

3.1 随机逼近

3.2 蒙特卡洛值估计

3.3 时序差分值估计

3.4 算法总结

# 第三章 值估计

## 3.1 随机逼近

## 3.2 蒙特卡洛值估计

## 3.3 时序差分值估计

## 3.4 算法总结

# 随机逼近

## 累计平均逼近 $E(x)$

*For time  $k$ :*

$$S \leftarrow S + x_k;$$

$$N \leftarrow N + 1;$$

$$u_k \leftarrow S/N;$$

$$N \rightarrow \infty, \quad u \rightarrow E(x)$$

# 随机逼近

## 增量均值逼近 $E(x)$

For time  $k$ :

$$N \leftarrow N+1;$$

$$u_k \leftarrow u_{k-1} + (1/N) (x_k - u_{k-1});$$

相当于均值的更新：

$$\begin{aligned}\mu_k &= \frac{1}{k} \sum_{j=1}^k x_j \\ &= \frac{1}{k} \left( x_k + \sum_{j=1}^{k-1} x_j \right) \\ &= \frac{1}{k} (x_k + (k-1)\mu_{k-1}) \\ &= \mu_{k-1} + \frac{1}{k} (x_k - \mu_{k-1})\end{aligned}$$

# 随机逼近

## Robbins-Monro 逼近 $E(x)$

For time  $k$ :

$$N \leftarrow N+1;$$

$$u_k \leftarrow u_{k-1} + a(x_k - u_{k-1});$$

相当于权重比例更新：

$$u_k \leftarrow (1 - a)u_{k-1} + ax_k$$

## 本课程常用 Robbins-Monro 随机逼近公式

# 第三章 值估计

3.1 随机逼近

**3.2 蒙特卡洛值估计**

3.3 时序差分值估计

3.4 算法总结

# 蒙特卡洛值估计

## 问题描述

- Goal: learn  $v_\pi$  from episodes of experience under policy  $\pi$

$$S_1, A_1, R_2, \dots, S_k \sim \pi$$

- Recall that the *return* is the total discounted reward:

$$G_t = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} \underline{R_T}$$

- Recall that the value function is the expected return:

$$v_\pi(s) = \mathbb{E}_\pi [G_t \mid S_t = s]$$

- Monte-Carlo policy evaluation uses *empirical mean* return instead of *expected* return



# 蒙特卡洛值估计

## First-Visit MC Evaluation

- To evaluate state  $s$
- The **first** time-step  $t$  that state  $s$  is visited in an episode,
- Increment counter  $N(s) \leftarrow N(s) + 1$
- Increment total return  $S(s) \leftarrow S(s) + G_t$
- Value is estimated by mean return  $V(s) = S(s)/N(s)$
- By law of large numbers,  $V(s) \rightarrow v_\pi(s)$  as  $N(s) \rightarrow \infty$

**只对 episode 的起始状态进行统计。**

# 蒙特卡洛值估计

## Every-Visit MC Evaluation

- To evaluate state  $s$
- **Every** time-step  $t$  that state  $s$  is visited in an episode,
- Increment counter  $N(s) \leftarrow N(s) + 1$
- Increment total return  $S(s) \leftarrow S(s) + G_t$
- Value is estimated by mean return  $V(s) = S(s)/N(s)$
- Again,  $V(s) \rightarrow v_\pi(s)$  as  $N(s) \rightarrow \infty$

**episode 的每个状态进行统计。**

## Incremental MC Evaluation

- Update  $V(s)$  incrementally after episode  $S_1, A_1, R_2, \dots, S_T$
- For each state  $S_t$  with return  $G_t$

$$N(S_t) \leftarrow N(S_t) + 1$$

$$V(S_t) \leftarrow V(S_t) + \frac{1}{N(S_t)} (G_t - V(S_t))$$

## Incremental MC Evaluation

- Update  $V(s)$  incrementally after episode  $S_1, A_1, R_2, \dots, S_T$
- For each state  $S_t$  with return  $G_t$

$$N(S_t) \leftarrow N(S_t) + 1$$

$$V(S_t) \leftarrow V(S_t) + \frac{1}{N(S_t)} (G_t - V(S_t))$$

- In non-stationary problems, it can be useful to track a running mean, i.e. forget old episodes.

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t - V(S_t))$$

# 第三章 值估计

3.1 随机逼近

3.2 蒙特卡洛值估计

**3.3 时序差分值估计**

3.4 算法总结

# 时序差分值估计

## 问题描述

采用**不完整的 episodes**, 估计  $V$  值。

- TD methods learn directly from episodes of experience
- TD is *model-free*: no knowledge of MDP transitions / rewards
- TD learns from *incomplete* episodes, by *bootstrapping*
- TD updates a guess towards a guess

# 时序差分值估计

## TD (0)

- Update value  $V(S_t)$  toward *estimated* return  $R_{t+1} + \gamma V(S_{t+1})$

$$V(S_t) \leftarrow V(S_t) + \alpha (R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$

Bellman 迭代的随机形式

- $R_{t+1} + \gamma V(S_{t+1})$  is called the *TD target*
- $\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$  is called the *TD error*

# 时序差分值估计

## TD (0)

### ■ V.S. MC

- Incremental every-visit Monte-Carlo
  - Update value  $V(S_t)$  toward *actual* return  $G_t$

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t - V(S_t))$$



# 时序差分值估计

## TD (0)

- TD can learn *before* knowing the final outcome
  - TD can learn online after every step
  - MC must wait until end of episode before return is known
- TD can learn *without* the final outcome
  - TD can learn from incomplete sequences
  - MC can only learn from complete sequences
  - TD works in continuing (non-terminating) environments
  - MC only works for episodic (terminating) environments

# 时序差分值估计

## TD (0)

- Return  $G_t = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T$  is *unbiased* estimate of  $v_\pi(S_t)$
- True TD target  $R_{t+1} + \gamma v_\pi(S_{t+1})$  is *unbiased* estimate of  $v_\pi(S_t)$
- TD target  $R_{t+1} + \gamma V(S_{t+1})$  is *biased* estimate of  $v_\pi(S_t)$
- TD target is much lower variance than the return:
  - Return depends on *many* random actions, transitions, rewards
  - TD target depends on *one* random action, transition, reward

# 时序差分值估计

## TD (0)

- MC has high variance, zero bias
  - Good convergence properties
  - (even with function approximation)
  - Not very sensitive to initial value
  - Very simple to understand and use
- TD has low variance, some bias
  - Usually more efficient than MC
  - TD(0) converges to  $v_{\pi}(s)$
  - (but not always with function approximation)
  - More sensitive to initial value

# 时序差分值估计

## TD (0)

- TD exploits Markov property
  - Usually more efficient in Markov environments
- MC does not exploit Markov property
  - Usually more effective in non-Markov environments

# 时序差分值估计

## TD (0)

### ■ AB Example:

Two states  $A, B$ ; no discounting; 8 episodes of experience

$A, 0, B, 0$

$B, 1$

$B, 1$

$B, 1$

$B, 1$

$B, 1$

$B, 1$

$B, 0$

What is  $V(A), V(B)$ ?

# 时序差分值估计

## TD (0)

- MC converges to solution with minimum mean-squared error
  - Best fit to the observed returns

$$\sum_{k=1}^K \sum_{t=1}^{T_k} (G_t^k - V(s_t^k))^2$$

- In the AB example,  $V(A) = 0$
- TD(0) converges to solution of max likelihood Markov model
  - Solution to the MDP  $\langle \mathcal{S}, \mathcal{A}, \hat{\mathcal{P}}, \hat{\mathcal{R}}, \gamma \rangle$  that best fits the data

$$\hat{\mathcal{P}}_{s,s'}^a = \frac{1}{N(s,a)} \sum_{k=1}^K \sum_{t=1}^{T_k} \mathbf{1}(s_t^k, a_t^k, s_{t+1}^k = s, a, s')$$

$$\hat{\mathcal{R}}_s^a = \frac{1}{N(s,a)} \sum_{k=1}^K \sum_{t=1}^{T_k} \mathbf{1}(s_t^k, a_t^k = s, a) r_t^k$$


- In the AB example,  $V(A) = 0.75$

# 时序差分值估计

## TD (0)

### TD Backup

$$V(S_t) \leftarrow V(S_t) + \alpha (R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$




# 时序差分值估计

## TD (0)

### V.S. MC

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t - V(S_t))$$






# 时序差分值估计

## TD (0)


### V.S. DP

$$V(S_t) \leftarrow \mathbb{E}_{\pi} [R_{t+1} + \gamma V(S_{t+1})]$$



# 时序差分值估计

## TD (0)



# 时序差分值估计

## TD (0)

### ■ Bootstrapping and Sampling

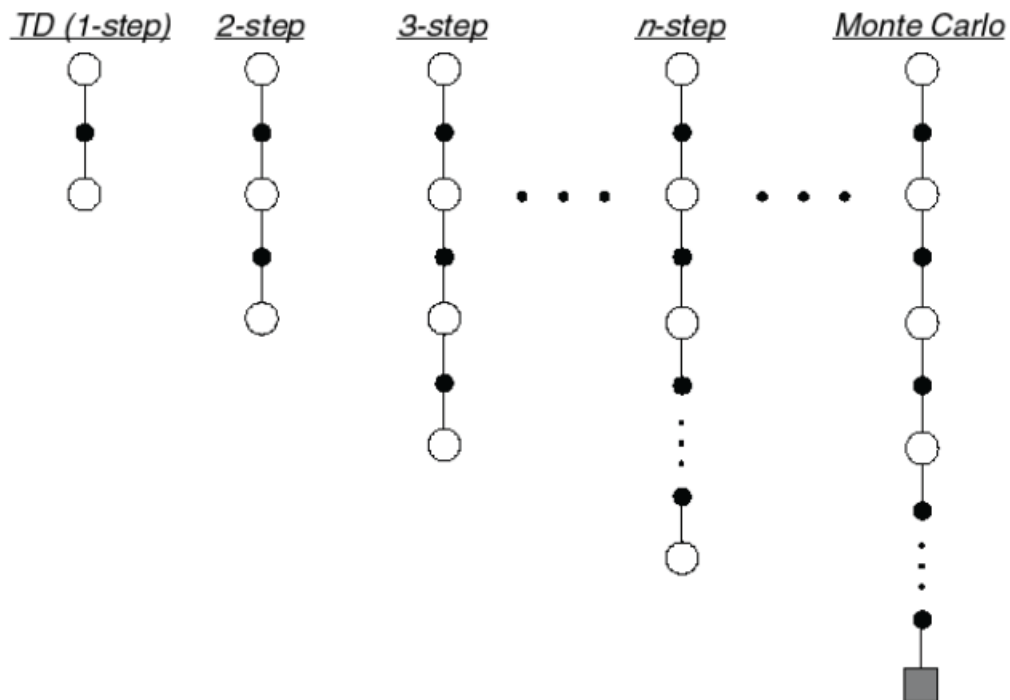
- **Bootstrapping**: update involves an estimate
  - MC does not bootstrap
  - DP bootstraps
  - TD bootstraps
- **Sampling**: update samples an expectation
  - MC samples
  - DP does not sample
  - TD samples

# 时序差分值估计

## TD( $\lambda$ )

### ■ n-step TD

- Let TD target look  $n$  steps into the future



# 时序差分值估计

## TD( $\lambda$ )

- Consider the following  $n$ -step returns for  $n = 1, 2, \infty$ :

$$\begin{array}{ll} n = 1 & (TD) \quad G_t^{(1)} = R_{t+1} + \gamma V(S_{t+1}) \\ n = 2 & \quad \quad G_t^{(2)} = R_{t+1} + \gamma R_{t+2} + \gamma^2 V(S_{t+2}) \\ & \quad \quad \vdots \\ n = \infty & (MC) \quad G_t^{(\infty)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T \end{array}$$

- Define the  $n$ -step return

$$G_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n V(S_{t+n})$$


- $n$ -step temporal-difference learning

$$V(S_t) \leftarrow V(S_t) + \alpha \left( G_t^{(n)} - V(S_t) \right)$$

# 时序差分值估计

## TD( $\lambda$ )

### ■ Forward Returns and TD( $\lambda$ )



- The  $\lambda$ -return  $G_t^\lambda$  combines all  $n$ -step returns  $G_t^{(n)}$
- Using weight  $(1-\lambda)\lambda^{n-1}$

$$G_t^\lambda = (1-\lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$$

- Forward-view TD( $\lambda$ )

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t^\lambda - V(S_t))$$

# 时序差分值估计

## TD( $\lambda$ )


$$\alpha \left( G_k^\lambda - V(S_k) \right) = \alpha \sum_{t=k}^T (\gamma \lambda)^{t-k} \delta_t$$

$$\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$$

$$V(s) \leftarrow V(s) + \alpha \sum_{t=k}^T (\gamma \lambda)^{t-k} \delta_t$$

# 时序差分值估计

## TD( $\lambda$ )




$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$$



# 时序差分值估计

## TD( $\lambda$ )



- Update value function towards the  $\lambda$ -return
- Forward-view looks into the future to compute  $G_t^\lambda$
- Like MC, can only be computed from complete episodes

# 时序差分值估计

## TD( $\lambda$ )

### ■ Backward TD( $\lambda$ )

$$E_0(s) = 0$$

$$E_t(s) = \gamma\lambda E_{t-1}(s) + \mathbf{1}(S_t = s)$$


- Keep an eligibility trace for every state  $s$
- Update value  $V(s)$  for every state  $s$
- In proportion to TD-error  $\delta_t$  and eligibility trace  $E_t(s)$

$$\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$$

$$V(s) \leftarrow V(s) + \alpha \delta_t E_t(s)$$

# 时序差分值估计

## TD( $\lambda$ )



# 时序差分值估计

## TD( $\lambda$ )

### ■ Forward and Backward TD

#### Theorem

*The sum of offline updates is identical for forward-view and backward-view TD( $\lambda$ )*

$$\sum_{t=1}^T \alpha \delta_t E_t(s) = \sum_{t=1}^T \alpha \left( G_t^\lambda - V(S_t) \right) \mathbf{1}(S_t = s)$$

# 时序差分值估计

## TD( $\lambda$ )

### ■ TD(0), TD( $\lambda$ ), TD(1)

#### TD(0)

$$E_t(s) = \mathbf{1}(S_t = s)$$

$$V(s) \leftarrow V(s) + \alpha \delta_t E_t(s)$$

This is exactly equivalent to TD(0) update

$$V(S_t) \leftarrow V(S_t) + \alpha \delta_t$$

# 时序差分值估计

## TD( $\lambda$ )

### TD( $\lambda$ )

$$E_t(s) = \gamma \lambda E_{t-1}(s) + \mathbf{1}(S_t = s)$$

$$V(s) \leftarrow V(s) + \alpha \delta_t E_t(s)$$

### TD(1)等价于MC

$$\alpha \left( G_k^\lambda - V(S_k) \right) = \alpha \sum_{t=k}^T (\gamma \lambda)^{t-k} \delta_t = \alpha \sum_{t=k}^T \gamma^{t-k} \delta_t = \alpha (G_k - V(S_k))$$

大家请自行证明

# 第三章 值估计

3.1 随机逼近

3.2 蒙特卡洛值估计

3.3 时序差分值估计

**3.4 算法总结**

# 算法总结

Offline updates	$\lambda = 0$	$\lambda \in (0, 1)$	$\lambda = 1$
Backward view	TD(0) 	TD( $\lambda$ ) 	TD(1) 
Forward view	TD(0)	Forward TD( $\lambda$ )	MC
Online updates	$\lambda = 0$	$\lambda \in (0, 1)$	$\lambda = 1$
Backward view	TD(0) 	TD( $\lambda$ ) <del>  </del>	TD(1) <del>  </del>
Forward view	TD(0) 	Forward TD( $\lambda$ ) 	MC 
Exact Online	TD(0)	Exact Online TD( $\lambda$ )	Exact Online TD(1)



# 本讲参考文献

1. Richard S. Sutton and Andrew G. Barto. Reinforcement Learning: An Introduction. (Second edition, in progress, draft.
2. David Silver, Slides@ 《Reinforcement Learning: An Introduction》, 2016.
3. Simon Haykin, 申富饶等译, 神经网络与学习机器, 第三版。